

# Aggregation

MongoDB provides a number of options for aggregating data. As usual we will start this lesson by inserting some example data.

```
In [1]: import pymongo
conn = pymongo.Connection()
db = conn.examples
db.things.insert({'x': 1, 'tags': ['dog', 'cat']})
db.things.insert({'x': 2, 'tags': ['cat']})
db.things.insert({'x': 3, 'tags': ['mouse', 'cat', 'dog']})
db.things.insert({'x': 4, 'tags': []})
```

```
Out[1]: ObjectId('4f56d13afba5224d6b000003')
```

The simplest aggregation method is `count()`

```
In [2]: db.things.count()
```

```
Out[2]: 4
```

```
In [3]: db.things.find({'x': 2}).count()
```

```
Out[3]: 1
```

MongoDB executes javascript server side for more advanced aggregation operations. Here's an example of using PyMongo's `map_reduce` method. The javascript map and reduce functions are defined using instances of `bson.code.Code`.

```
In [4]: from bson.code import Code
mymap = Code("function () {"
            "  this.tags.forEach(function(z) {"
            "    emit(z, 1);"
            "  });"
            "}")
myreduce = Code("function (key, values) {"
               "  var total = 0;"
               "  for (var i = 0; i < values.length; i++) {"
               "    total += values[i];"
               "  }"
               "  return total;"
               "}")
coll = db.things.map_reduce(mymap, myreduce, "myresults")
for doc in coll.find(): print doc

{'_id': u'cat', 'value': 3.0}
{'_id': u'dog', 'value': 2.0}
{'_id': u'mouse', 'value': 1.0}
```

The output of `map_reduce` is stored in the collection "myresults". If we didn't want to store the results we could use PyMongo's `inline_map_reduce` method instead. The results would be returned in a list.

PyMongo also provides a `group()` method for doing group operations with javascript. Group will be covered in the exercises at the end of this lesson.

## GridFS

A single MongoDB document is limited to 16MB in size. This is generally large enough for textual data but what if you want to store large binary files in MongoDB? GridFS is the answer. GridFS is a protocol implemented by PyMongo to store binary data in document "chunks" on the server, bypassing the document size limit.

Here's a simple example inserting some text using the `gridfs` module.

```
In [5]: import gridfs
        db = conn.gridfs_example
        gfs = gridfs.GridFS(db)
        a = gfs.put("Hello PyCon!")
```

Now lets read the data back. GridFS is implemented using file-like objects. The `get` method returns a `GridOut` object that provides file methods like `read`, `readline`, `seek`, `tell`, and `close`.

```
In [6]: f = gfs.get(a)
        f.read()
```

```
Out[6]: 'Hello PyCon!'
```

Here's another example inserting the same file but including some metadata to be stored with it.

```
In [7]: b = gfs.put(gfs.get(a), filename="foo", bar="baz")
        out = gfs.get(b)
        out.read()
```

```
Out[7]: 'Hello PyCon!'
```

The file metadata can be accessed as attributes of the `GridOut` object.

```
In [8]: out.filename
```

```
Out[8]: u'foo'
```

```
In [9]: out.bar
```

```
Out[9]: u'baz'
```

An upload date is stored with each file in `gridfs`.

```
In [10]: out.upload_date
```

```
Out[10]: datetime.datetime(2012, 3, 7, 3, 8, 48, 230000)
```

## Exercises

- I. Implement the group example from <http://www.mongodb.org/display/DOCS/Aggregation#Aggregation-UsingGroupfromVariousLanguages>
- II. Store a file from your home directory in MongoDB using GridFS. Now read it back.